

AI Trader — Full Project Overview

1. Executive Context

AI Trader is a cloud-native, autonomous equities trading platform designed to ingest market data deterministically, generate probabilistic signals, apply explicit risk controls, and execute trades in a fully observable and auditable way. The system prioritizes **risk-first design, traceability, and operational discipline** over speculative AI behavior.

The goal is not to "beat the market" indiscriminately, but to build a **repeatable, governable trading engine** that can evolve safely over time.

2. Guiding Principles

- Deterministic data ingestion (replayable, auditable)
 - Probabilistic signals (uncertainty-aware, not point forecasts)
 - Explicit risk gates (fractional Kelly, exposure caps, drawdown halts)
 - Shared pipelines for backtest and live trading
 - Cloud-native operations with first-class observability
-

3. High-Level Architecture

The system is organized into **phased layers**, each with a clear responsibility:

1. **Data & Sessions** – multi-source market data ingestion, session tagging
2. **Probabilistic Core** – Kalman filtering, regime detection
3. **Strategies & Backtesting** – signal generation and historical evaluation
4. **Risk & Execution** – sizing, gating, broker routing
5. **Observability & Ops** – logs, metrics, traces, dashboards
6. **CI/CD & Automation** – build, deploy, scheduling

This separation prevents tight coupling and allows safe iteration.

4. Codebase Structure

The codebase is modular and import-safe, designed for long-term maintainability.

```
app/
  ├── api/           # FastAPI routes, health checks
  └── dal/          # Market Data Abstraction Layer
```

```

├── probability/      # Kalman + regime pipeline
├── features/         # Indicators, MTF aggregation
├── scanners/         # Premarket & intraday scanners
├── services/          # Watchlist & orchestration services
├── strats/           # Strategy implementations
├── agent/             # Policy, sizing, risk, journaling
├── backtest/          # Backtest engine & broker simulator
├── execution/         # Live broker adapters
├── sessions/          # Market session calendar
├── observability/    # OTEL + logging helpers
├── monitoring/        # Streamlit dashboards
├── db/                # SQLAlchemy models & repositories
└── tests/             # Unit, integration, regression tests

```

Key design decisions:

- Strategies never talk directly to vendors
- Live and backtest share the same execution logic
- Probabilistic signals are first-class data artifacts

5. Market Data & Probabilistic Core

Market Data Abstraction Layer (DAL)

The DAL unifies multiple vendors (Alpaca, Alpha Vantage, Finnhub, Yahoo fallback) behind a single interface.

Responsibilities:

- Normalize OHLCV data into deterministic schemas
- Handle HTTP + streaming ingestion
- Perform automatic gap backfills
- Cache results to Parquet (Blob Storage)
- Optionally persist metadata to Postgres

Probabilistic Pipeline

- Kalman Filter produces smoothed price, velocity, and uncertainty
- Regime classifiers (HMM/GARCH planned) label volatility states
- Outputs feed strategy selection and risk sizing

This avoids overconfident signals and supports adaptive risk management.

6. Strategies, Risk & Execution

Strategies

- Breakout strategy: live and end-to-end wired
- Momentum / Mean-Reversion: scaffolded, pending parameter sweeps

Risk Management (Core Design)

- Fractional Kelly sizing (spec complete)
- Hard exposure caps (symbol + portfolio)
- Max drawdown halts
- Buying power and PDT guards
- Session-aware throttling

Execution

- Broker: Alpaca (paper/live)
- Order types:
- Limit / bracket orders in PRE/AFT
- Marketable limits during RTH
- OMS cache persisted to Postgres

Risk gates sit **before** any broker call.

7. Infrastructure (Azure)

Core Components

Layer	Technology
Compute	Azure App Service (Linux container)
Database	Azure PostgreSQL Flexible Server
Storage	Azure Blob Storage
Secrets	Azure Key Vault
Observability	OpenTelemetry + App Insights
CI/CD	GitHub Actions

Security & Identity

- Managed Identity everywhere possible
- No static credentials in code or repos
- Key Vault as the single source of truth

Database Design

Postgres schemas are domain-separated: - market - trading - backtest - analytics - storage

Backups, PITR, and restore drills are documented and tested.

8. Observability & Operations

Observability is a first-class concern:

- Structured JSON logging
- Distributed traces via OpenTelemetry
- Metrics for latency, errors, exposure, PnL
- Streamlit dashboards for live visibility
- Alerting for failures and risk breaches

Operational runbooks define:
- Daily trading schedule
- Data degradation behavior
- Broker failure handling
- Safe mode and recovery steps

9. CI/CD & Automation

Continuous Integration

- Linting (ruff)
- Formatting (black)
- Security scans (bandit, pip-audit)
- Pytest on every PR

Continuous Deployment

- GitHub Actions build containers
- Deploy to Azure App Service
- Slot-based promotion

Scheduling

- Webhook-driven jobs for:
 - Premarket scans
 - Watchlist refreshes
 - End-of-day processing
 - Model retraining
-

10. Project Status (Current)

Phase	Status
Bootstrap	Complete
Data & Sessions	Complete
Probabilistic Stack	~70%

Phase	Status
Strategy Coverage	~40%
Risk & Execution	~20%
Dashboard & Ops	~55%
CI/CD & Scheduling	~25%
Continuous Learning	~5%
Hardening	~5%

Primary focus areas: - Finish RiskManagementAgent - Wire execution loop fully - Enforce paper-trading discipline

11. Blunt Assessment

This is not a toy project.

Architecturally, the system is **sound, modern, and defensible**. The remaining risk is not technical—it is execution discipline and scope control.

If risk and execution are completed cleanly and paper trading is enforced rigorously, AI Trader is well-positioned to achieve its stated goal: a sustainable, autonomous trading engine with controlled downside and continuous learning potential.

Document purpose: single-source project context for future development, audits, or collaborators.